

# SELLING PRICE PREDICTION OF OLD VEHICLES

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: dataset=pd.read_csv("car_prediction_dataset.csv")
```

```
In [3]: dataset.shape
```

```
Out[3]: (301, 9)
```

```
In [4]: dataset.head(2)
```

```
Out[4]:
```

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission	owner
0	ritz	2014	5.59	27000	Petrol	Dealer	Manual	0
1	sx4	2013	9.54	43000	Diesel	Dealer	Manual	0



```
In [5]: dataset.isnull().sum()
```

```
Out[5]: car_name      0
year      0
present_price  0
kms_driven  0
fuel_type   0
seller_type  0
transmission  0
owner      0
selling_price  0
dtype: int64
```

```
In [6]: dataset.duplicated().sum()
```

```
Out[6]: np.int64(2)
```

```
In [7]: dataset.drop_duplicates()
```

```
Out[7]:
```

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission	owne
0	ritz	2014	5.59	27000	Petrol	Dealer	Manual	
1	sx4	2013	9.54	43000	Diesel	Dealer	Manual	
2	ciaz	2017	9.85	6900	Petrol	Dealer	Manual	
3	wagon r	2011	4.15	5200	Petrol	Dealer	Manual	
4	swift	2014	6.87	42450	Diesel	Dealer	Manual	
...	...	...	...	...	...	...	...	...
296	city	2016	11.60	33988	Diesel	Dealer	Manual	
297	brio	2015	5.90	60000	Petrol	Dealer	Manual	
298	city	2009	11.00	87934	Petrol	Dealer	Manual	
299	city	2017	12.50	9000	Diesel	Dealer	Manual	
300	brio	2016	5.90	5464	Petrol	Dealer	Manual	

299 rows × 9 columns



```
In [8]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   car_name        301 non-null    object
1   year            301 non-null    int64
2   present_price   301 non-null    float64
3   kms_driven      301 non-null    int64
4   fuel_type       301 non-null    object
5   seller_type     301 non-null    object
6   transmission    301 non-null    object
7   owner           301 non-null    int64
8   selling_price   301 non-null    float64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [9]: dataset.describe()
```

```
Out[9]:
```

	year	present_price	kms_driven	owner	selling_price
<b>count</b>	301.000000	301.000000	301.000000	301.000000	301.000000
<b>mean</b>	2013.627907	7.628472	36947.205980	0.043189	4.661296
<b>std</b>	2.891554	8.644115	38886.883882	0.247915	5.082812
<b>min</b>	2003.000000	0.320000	500.000000	0.000000	0.100000
<b>25%</b>	2012.000000	1.200000	15000.000000	0.000000	0.900000
<b>50%</b>	2014.000000	6.400000	32000.000000	0.000000	3.600000
<b>75%</b>	2016.000000	9.900000	48767.000000	0.000000	6.000000
<b>max</b>	2018.000000	92.600000	500000.000000	3.000000	35.000000

## LABEL ENCODING

```
In [10]: from sklearn.preprocessing import LabelEncoder
```

### CAR NAME ENCODING

```
In [11]: car_name_le=LabelEncoder()
dataset['car_name']=car_name_le.fit_transform(dataset['car_name'])
```

### FUEL TYPE ENCODING

```
In [12]: dataset['fuel_type'].unique()
```

```
Out[12]: array(['Petrol', 'Diesel', 'CNG'], dtype=object)
```

```
In [13]: fuel_type_le=LabelEncoder()
dataset['fuel_type']=fuel_type_le.fit_transform(dataset['fuel_type'])
```

### SELLER TYPE ENCODING

```
In [14]: seller_type_le=LabelEncoder()
dataset['seller_type']=seller_type_le.fit_transform(dataset['seller_type'])
```

### TRANSMISSION ENCODING

```
In [15]: transmission_le=LabelEncoder()
dataset['transmission']=transmission_le.fit_transform(dataset['transmission'])
```

## VISUALISING THE DATA


```
In [16]: dataset.head(3)
```

```
Out[16]:
```

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission	owner
0	90	2014	5.59	27000	2	0	1	0
1	93	2013	9.54	43000	1	0	1	0
2	68	2017	9.85	6900	2	0	1	0



```
In [17]: sns.heatmap(data=dataset.corr(),annot=True)
plt.show()
```



```
In [18]: input_data=dataset.iloc[:, :-1]
output_data=dataset['selling_price']
```

## SCALING DATA TO REMOVE ANAMOLY


```
In [19]: from sklearn.preprocessing import StandardScaler
```

```
In [20]: ss=StandardScaler()
input_data=pd.DataFrame(ss.fit_transform(input_data),columns=input_data.columns)
```

```
In [21]: dataset.head(3)
```

Out[21]:

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission	owner
0	90	2014	5.59	27000	2	0	1	0
1	93	2013	9.54	43000	1	0	1	0
2	68	2017	9.85	6900	2	0	1	0



## SPLITTING DATA

In [22]: `from sklearn.model_selection import train_test_split`

In [23]: `x_train,x_test,y_train,y_test=train_test_split(input_data,output_data,test_size=0.2`

## MODEL SELECTION

In [24]: `from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet  
from sklearn.tree import DecisionTreeRegressor  
from sklearn.svm import SVR  
from sklearn.neighbors import KNeighborsRegressor  
from sklearn.ensemble import RandomForestRegressor`

In [25]: `from sklearn.metrics import mean_squared_error,mean_absolute_error`

In [26]: `lr=LinearRegression()  
lr.fit(x_train,y_train)  
lr.score(x_train,y_train)*100, lr.score(x_test,y_test)*100`

Out[26]: (88.40630578239454, 84.65539666857805)

In [27]: `mean_squared_error(y_test,lr.predict(x_test)),mean_absolute_error(y_test,lr.predict`

Out[27]: (3.5347190697129753, 1.2142339415047014)

In [28]: `lr1=Lasso(alpha=0.5)  
lr1.fit(x_train,y_train)  
lr1.score(x_train,y_train)*100, lr1.score(x_test,y_test)*100`

Out[28]: (85.0124395411389, 78.63348352385346)

In [29]: `mean_squared_error(y_test,lr1.predict(x_test)),mean_absolute_error(y_test,lr1.predi`

Out[29]: (4.921901961904479, 1.4658722139390727)

In [30]: `lr2=Ridge(alpha=10)  
lr2.fit(x_train,y_train)  
lr2.score(x_train,y_train)*100, lr2.score(x_test,y_test)*100`

Out[30]: (88.28628537091497, 84.16213595432284)

```
In [31]: mean_squared_error(y_test,lr2.predict(x_test)),mean_absolute_error(y_test,lr2.predict(x_test))
```

```
Out[31]: (3.6483445584506238, 1.233161689725921)
```

```
In [32]: lr3=ElasticNet(alpha=0.5)
lr3.fit(x_train,y_train)
lr3.score(x_train,y_train)*100, lr3.score(x_test,y_test)*100
```

```
Out[32]: (84.00059239671332, 78.3177718663528)
```

```
In [33]: mean_squared_error(y_test,lr3.predict(x_test)),mean_absolute_error(y_test,lr3.predict(x_test))
```

```
Out[33]: (4.99462798760845, 1.4094819518715052)
```

```
In [34]: dt=DecisionTreeRegressor(max_depth=15)
dt.fit(x_train,y_train)
dt.score(x_train,y_train)*100, dt.score(x_test,y_test)*100
```

```
Out[34]: (100.0, 94.52366873121181)
```

```
In [35]: mean_squared_error(y_test,dt.predict(x_test)),mean_absolute_error(y_test,dt.predict(x_test))
```

```
Out[35]: (1.2615049180327873, 0.6940983606557376)
```

```
In [36]: rf=RandomForestRegressor(n_estimators=100)
rf.fit(x_train,y_train)
rf.score(x_train,y_train)*100, rf.score(x_test,y_test)*100
```

```
Out[36]: (98.33237614607282, 96.59775818440349)
```

```
In [37]: mean_squared_error(y_test,rf.predict(x_test)),mean_absolute_error(y_test,rf.predict(x_test))
```

```
Out[37]: (0.7837262890163934, 0.5828016393442627)
```

```
In [38]: sv=SVR()
sv.fit(x_train,y_train)
sv.score(x_train,y_train)*100, sv.score(x_test,y_test)*100
```

```
Out[38]: (66.00840380338376, 78.48466914602926)
```

```
In [39]: mean_squared_error(y_test,sv.predict(x_test)),mean_absolute_error(y_test,sv.predict(x_test))
```

```
Out[39]: (4.956182223686514, 0.9945642205689335)
```

```
In [40]: knn=KNeighborsRegressor(n_neighbors=10)
knn.fit(x_train,y_train)
knn.score(x_train,y_train)*100, knn.score(x_test,y_test)*100
```

```
Out[40]: (86.59124637342433, 90.59004346126326)
```

```
In [41]: mean_squared_error(y_test,knn.predict(x_test)),mean_absolute_error(y_test,knn.predict(x_test))
```

Out[41]: (2.1676384918032783, 0.9017704918032786)

In [42]: *## now its clear that random forest is the best suitable model with good accuracy a*

In [43]: *x\_test #taking a test data to predict*

Out[43]:

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission
177	-1.275759	0.821718	-0.817924	-0.333500	0.500183	1.356327	-2.554408
289	0.251795	0.821718	0.691970	-0.668875	0.500183	-0.737285	0.391480
228	1.230996	-0.563924	0.205282	0.593804	-1.852241	-0.737285	0.391480
198	-2.059120	-0.910335	-0.817924	-0.050157	0.500183	1.356327	0.391480
60	0.330131	-0.217514	1.272521	0.078661	0.500183	-0.737285	0.391480
...	...	...	...	...	...	...	...
234	0.760979	0.475308	-0.223468	-0.835995	0.500183	-0.737285	0.391480
296	0.251795	0.821718	0.460214	-0.076225	-1.852241	-0.737285	0.391480
281	0.251795	-2.642389	-0.003299	0.347965	0.500183	-0.737285	0.391480
285	0.956819	0.821718	0.100991	-0.563806	0.500183	-0.737285	-2.554408
182	-1.158255	-0.217514	-0.816765	-0.178949	0.500183	1.356327	0.391480

61 rows × 8 columns



In [44]: `rf.predict([[ -1.275759,0.821718, -0.817924, -0.333500,0.500183,1.356327, -2.554408, -0.`

C:\Users\Aryansh Pathak\anaconda3\Lib\site-packages\sklearn\utils\validation.py:273  
9: UserWarning: X does not have valid feature names, but RandomForestRegressor was fitted with feature names  
warnings.warn(

Out[44]: array([0.4365])

In [45]: *y\_test #comparing the predicted data and test data*

Out[45]:

177	0.35
289	10.11
228	4.95
198	0.15
60	6.95
...	
234	5.50
296	9.50
281	2.10
285	7.40
182	0.30

Name: selling\_price, Length: 61, dtype: float64

```
In [46]: # Let's do a prediction for this dataset the original output is 3.35
#         car_name      year  present_price  kms_driven  fuel_type  sel
#         ritz         2014          5.59        27000    Petrol
```

```
In [47]: new_data=pd.DataFrame([[ 'ritz',2014,5.59,27000,'Petrol','Dealer','Manual',0]],column
```

```
In [48]: new_data
```

```
Out[48]:
```

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission	owner
0	ritz	2014	5.59	27000	Petrol	Dealer	Manual	0

```
In [49]: new_data['car_name']=car_name_le.fit_transform(new_data['car_name'])
```

```
In [50]: new_data['fuel_type']=fuel_type_le.fit_transform(new_data['fuel_type'])
```

```
In [51]: new_data['seller_type']=seller_type_le.fit_transform(new_data['seller_type'])
```

```
In [52]: new_data['transmission']=transmission_le.fit_transform(new_data['transmission'])
```

```
In [53]: new_data=pd.DataFrame(ss.transform(new_data),columns=new_data.columns)
```

```
In [54]: new_data
```

```
Out[54]:
```

	car_name	year	present_price	kms_driven	fuel_type	seller_type	transmission	
0	-2.4508	0.128897	-0.236215	-0.256224	-4.204665	-0.737285	-2.554408	-0.1



```
In [55]: rf.predict(new_data)
```

```
Out[55]: array([3.3675])
```

```
In [57]: # our predicted value is 3.3675
```

# Thank you !